

Bias and Variance and Overfitting

5

A

CS 3244
Machine Learning



NUS | Computing

Recap from Week 04

Cost Function for Logistic Regression



For logistic regression,

$$L_{\text{train}}(\theta) = \frac{1}{m} \sum_{j=1}^m \ln(1 + \exp^{-y^{(j)} \theta^T x^{(j)}})$$

Iterative
Solution


Compare to linear regression:

$$L_{\text{train}}(\theta) = \frac{1}{m} \sum_{j=1}^m (\theta^T x^{(j)} - y^{(j)})^2$$

Closed-form
solution, but
iteration also
possible

NUS CS3234: Machine Learning

47




GD vs. SGD on $m = 10$




GD

SGD



10 steps




30 steps

NUS CS3234: Machine Learning

56

Inherently handles noisy data



NUS CS3234: Machine Learning

58

Forecast for Week 05



Learning Outcomes for this week:

- Understand the bias–variance tradeoff
- Understand why overfitting occurs: the role of model complexity and the sampled dataset
- Apply bias and variance decomposition in simple scenarios



Bias and Variance, Illustrated

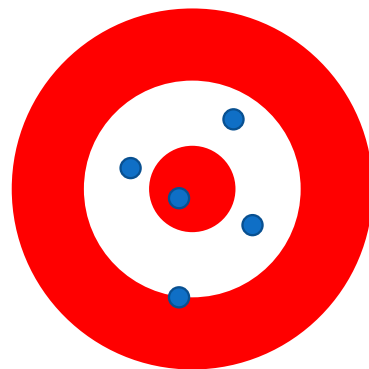
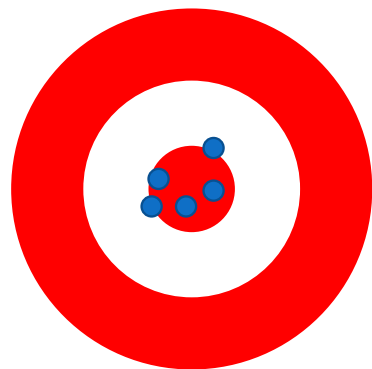
CS3244 Machine Learning



Department of Computer Science
School of Computing



Bias

The difference between the average prediction and the true value.



Variance

The variability of the model prediction, for given data.
Tells us spread of our data.



Regressing the sine function



$$f: [-1, 1] \rightarrow \mathbb{R}$$

Only two training examples! $m = 2$


Two models used for learning:

1 $\mathcal{H}_1: h(x) = \theta_1 x + \theta_0$

2 $\mathcal{H}_0: h(x) = \theta_0$



Your turn Q1: Which is better, \mathcal{H}_0 or \mathcal{H}_1 ?

$$f(x) = \sin(\pi x)$$





Approximation – \mathcal{H}_0 versus \mathcal{H}_1

Use the full power of the model





where — and — are the $\bar{h}(x)$ of each \mathcal{H} .


Learning – \mathcal{H}_0 versus \mathcal{H}_1



Bias & variance - \mathcal{H}_0 versus \mathcal{H}_1




And the winner is ...



$$\text{Bias} = 0.50$$

$$\text{Var} = ?$$



$$\text{Bias} = 0.21$$

$$\text{Var} = ?$$

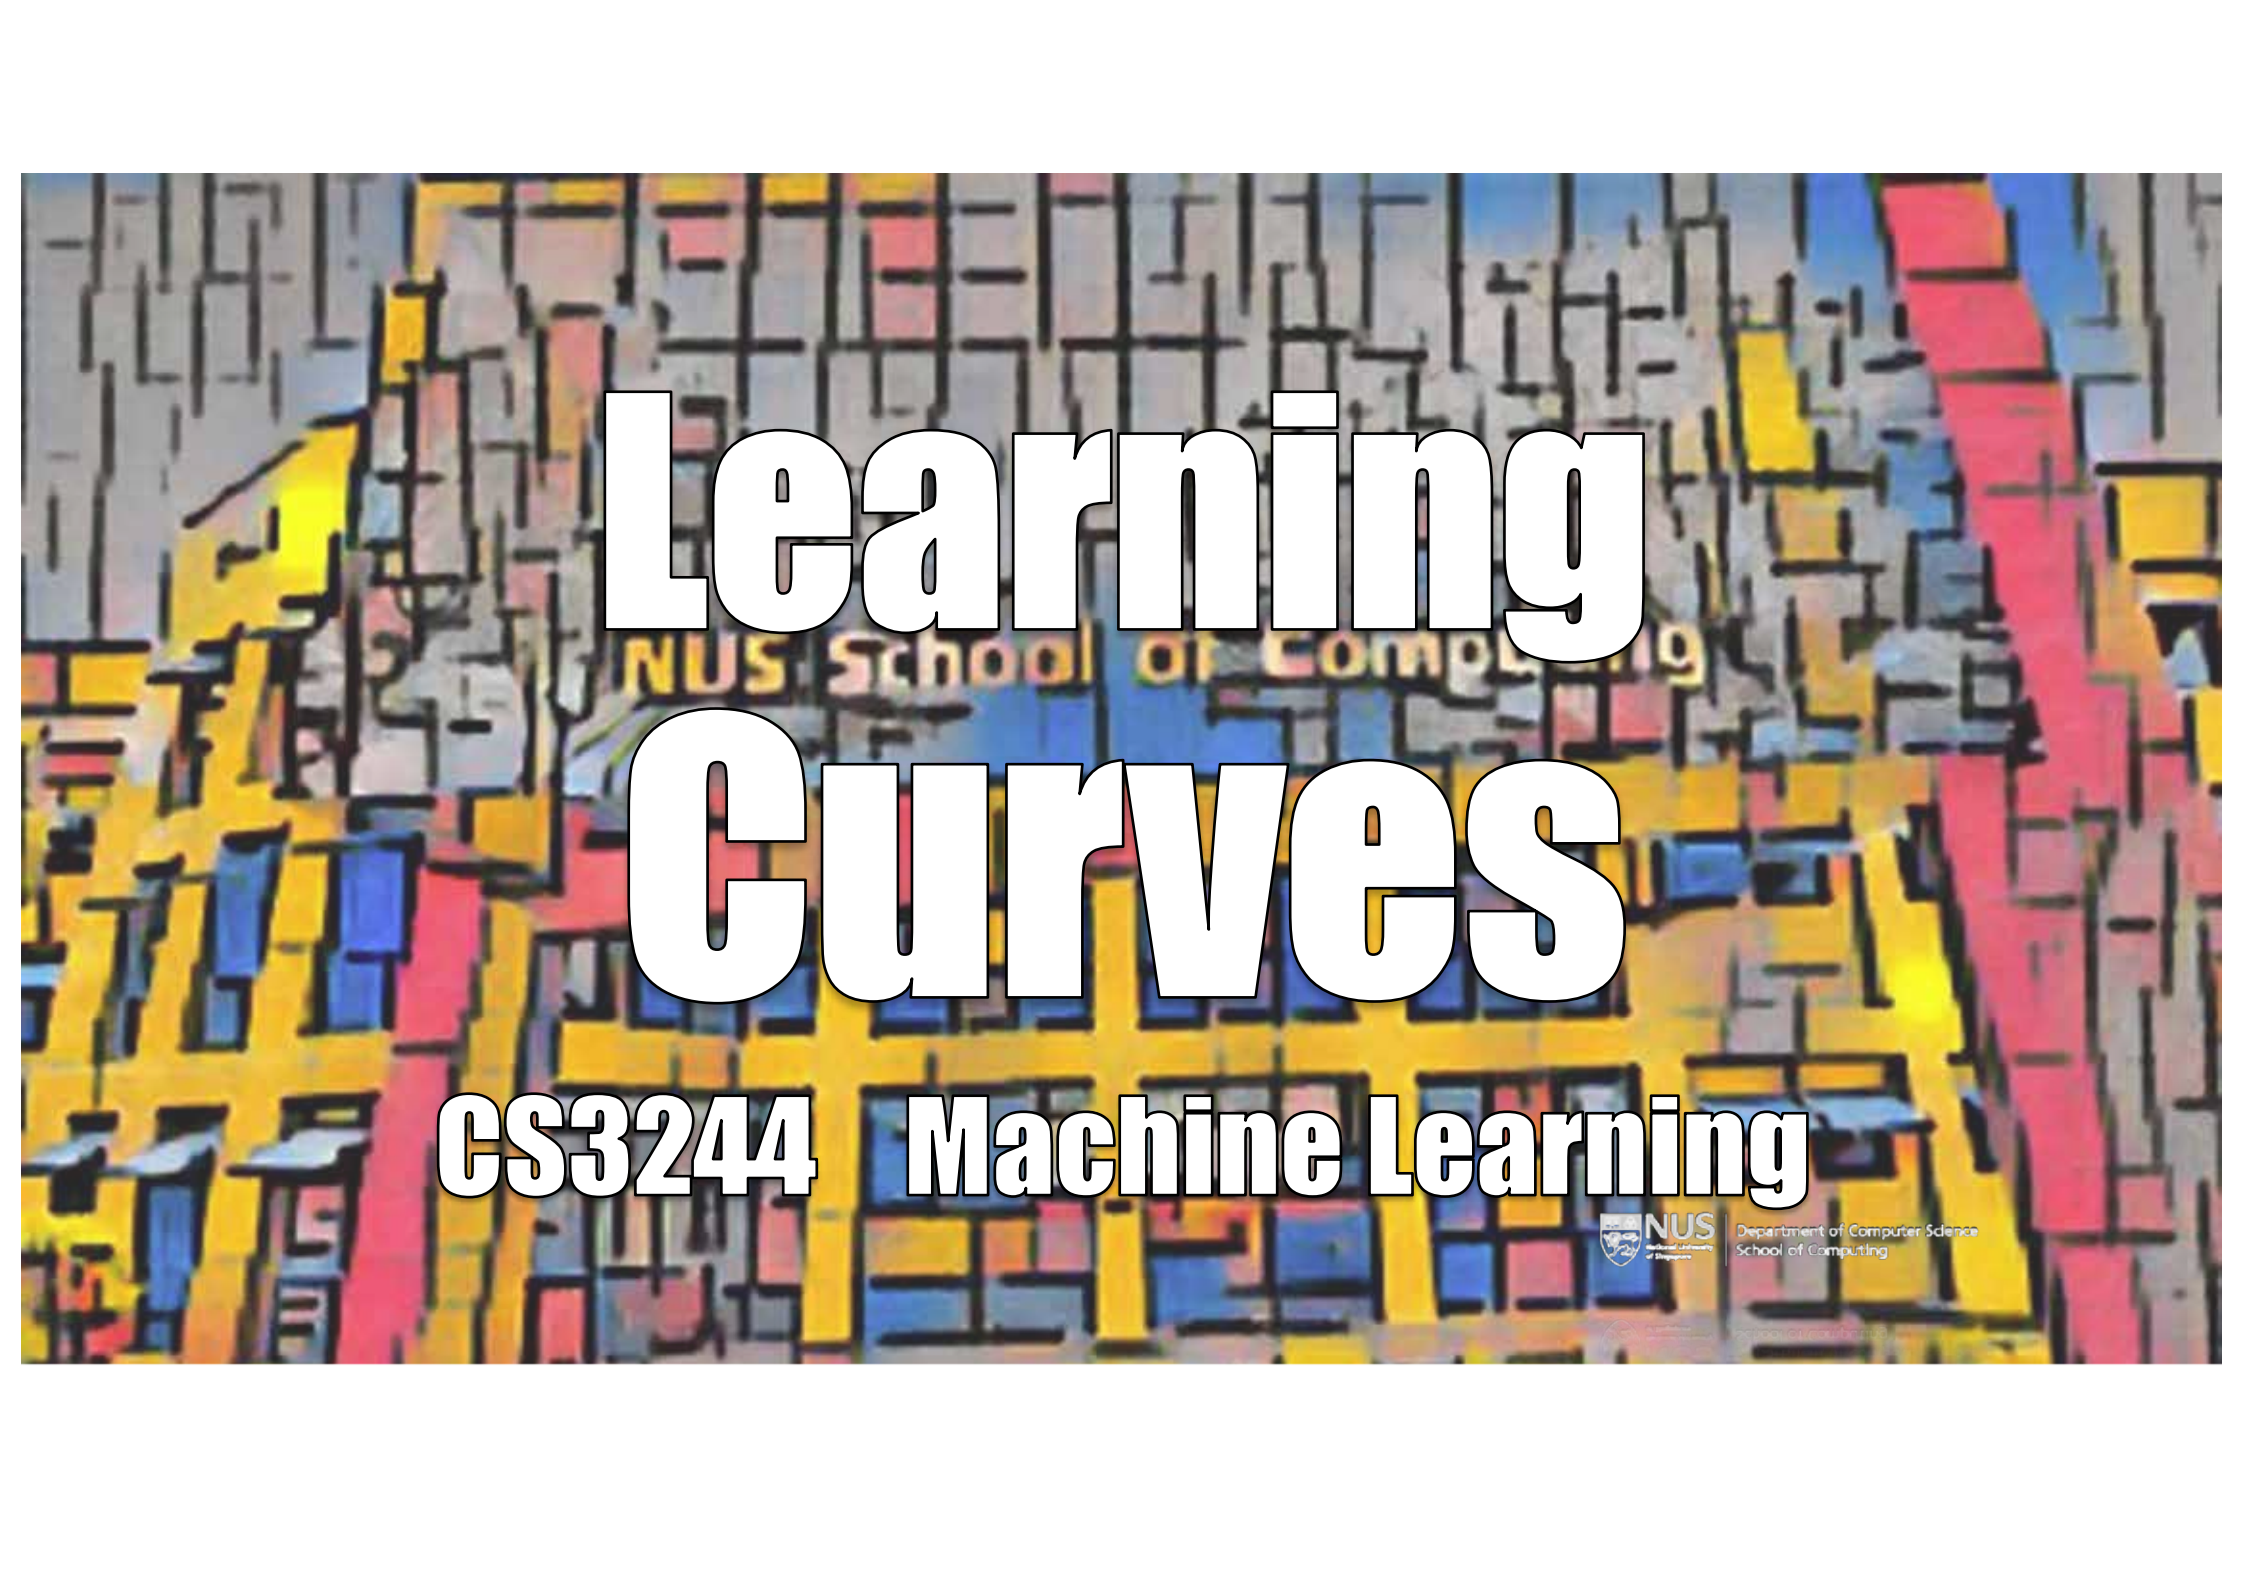
Lesson learned



Match the 'model complexity' to the

data resources, not to the **target complexity**.





Learning Curves

CS3244 Machine Learning



Department of Computer Science
School of Computing

Learning Curves

Plot expected L_{test} and L_{train} ...


...as we vary size m

Expected test cost $\mathbb{E}_{\mathcal{D}}[L_{test}(f_{\mathcal{D}})]$


Expected training cost $\mathbb{E}_{\mathcal{D}}[L_{train}(f_{\mathcal{D}})]$



Simple versus Complex Models

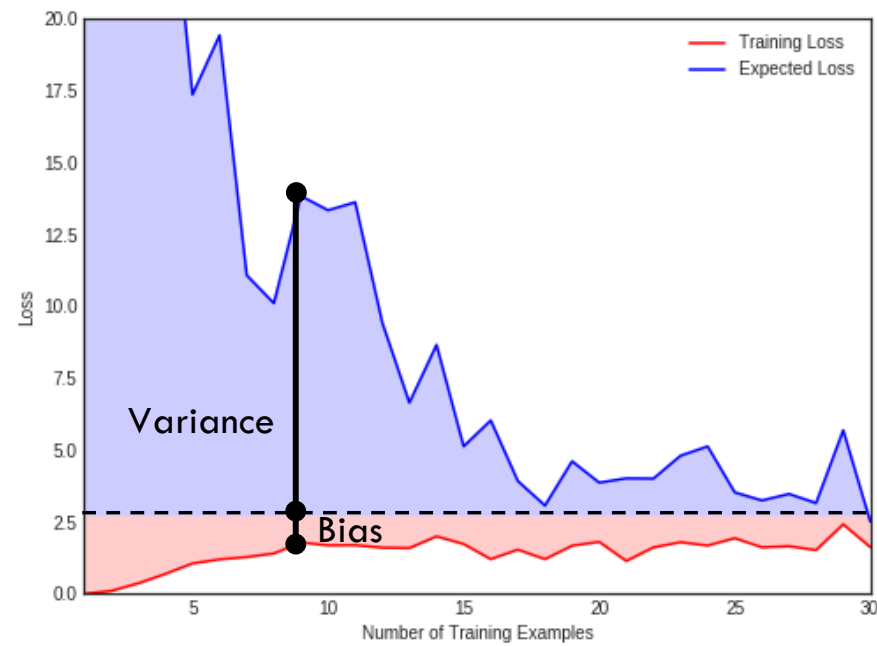



Simple Model



Complex Model

Bias–Variance on a Curve





NUS School of Computing

Overfitting

CS3244 Machine Learning



Department of Computer Science
School of Computing


Overfitting, Illustrated

Simple target function

5 data points –
slightly noisy

4th-order polynomial fit

$$L_{train} = 0$$



The culprit

Overfitting:

“Fitting the data more than is warranted”

Culprit: Fitting the noise.

Not just useless, **but harmful**


🥈 2nd versus 🥇 10th order polynomials



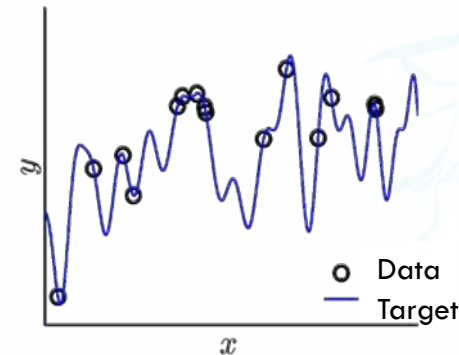
In Zoom breakout or physical subgroups, discuss and react 🏃🏃🏃:

(5 mins): Pick your model of choice for each scenario. In each scenario, you are given 15 data points. Justify your choice.

Q2: 10th-order target + noise




Q3: 50th order target, noiseless




Two fits for each function

Noisy low-order target

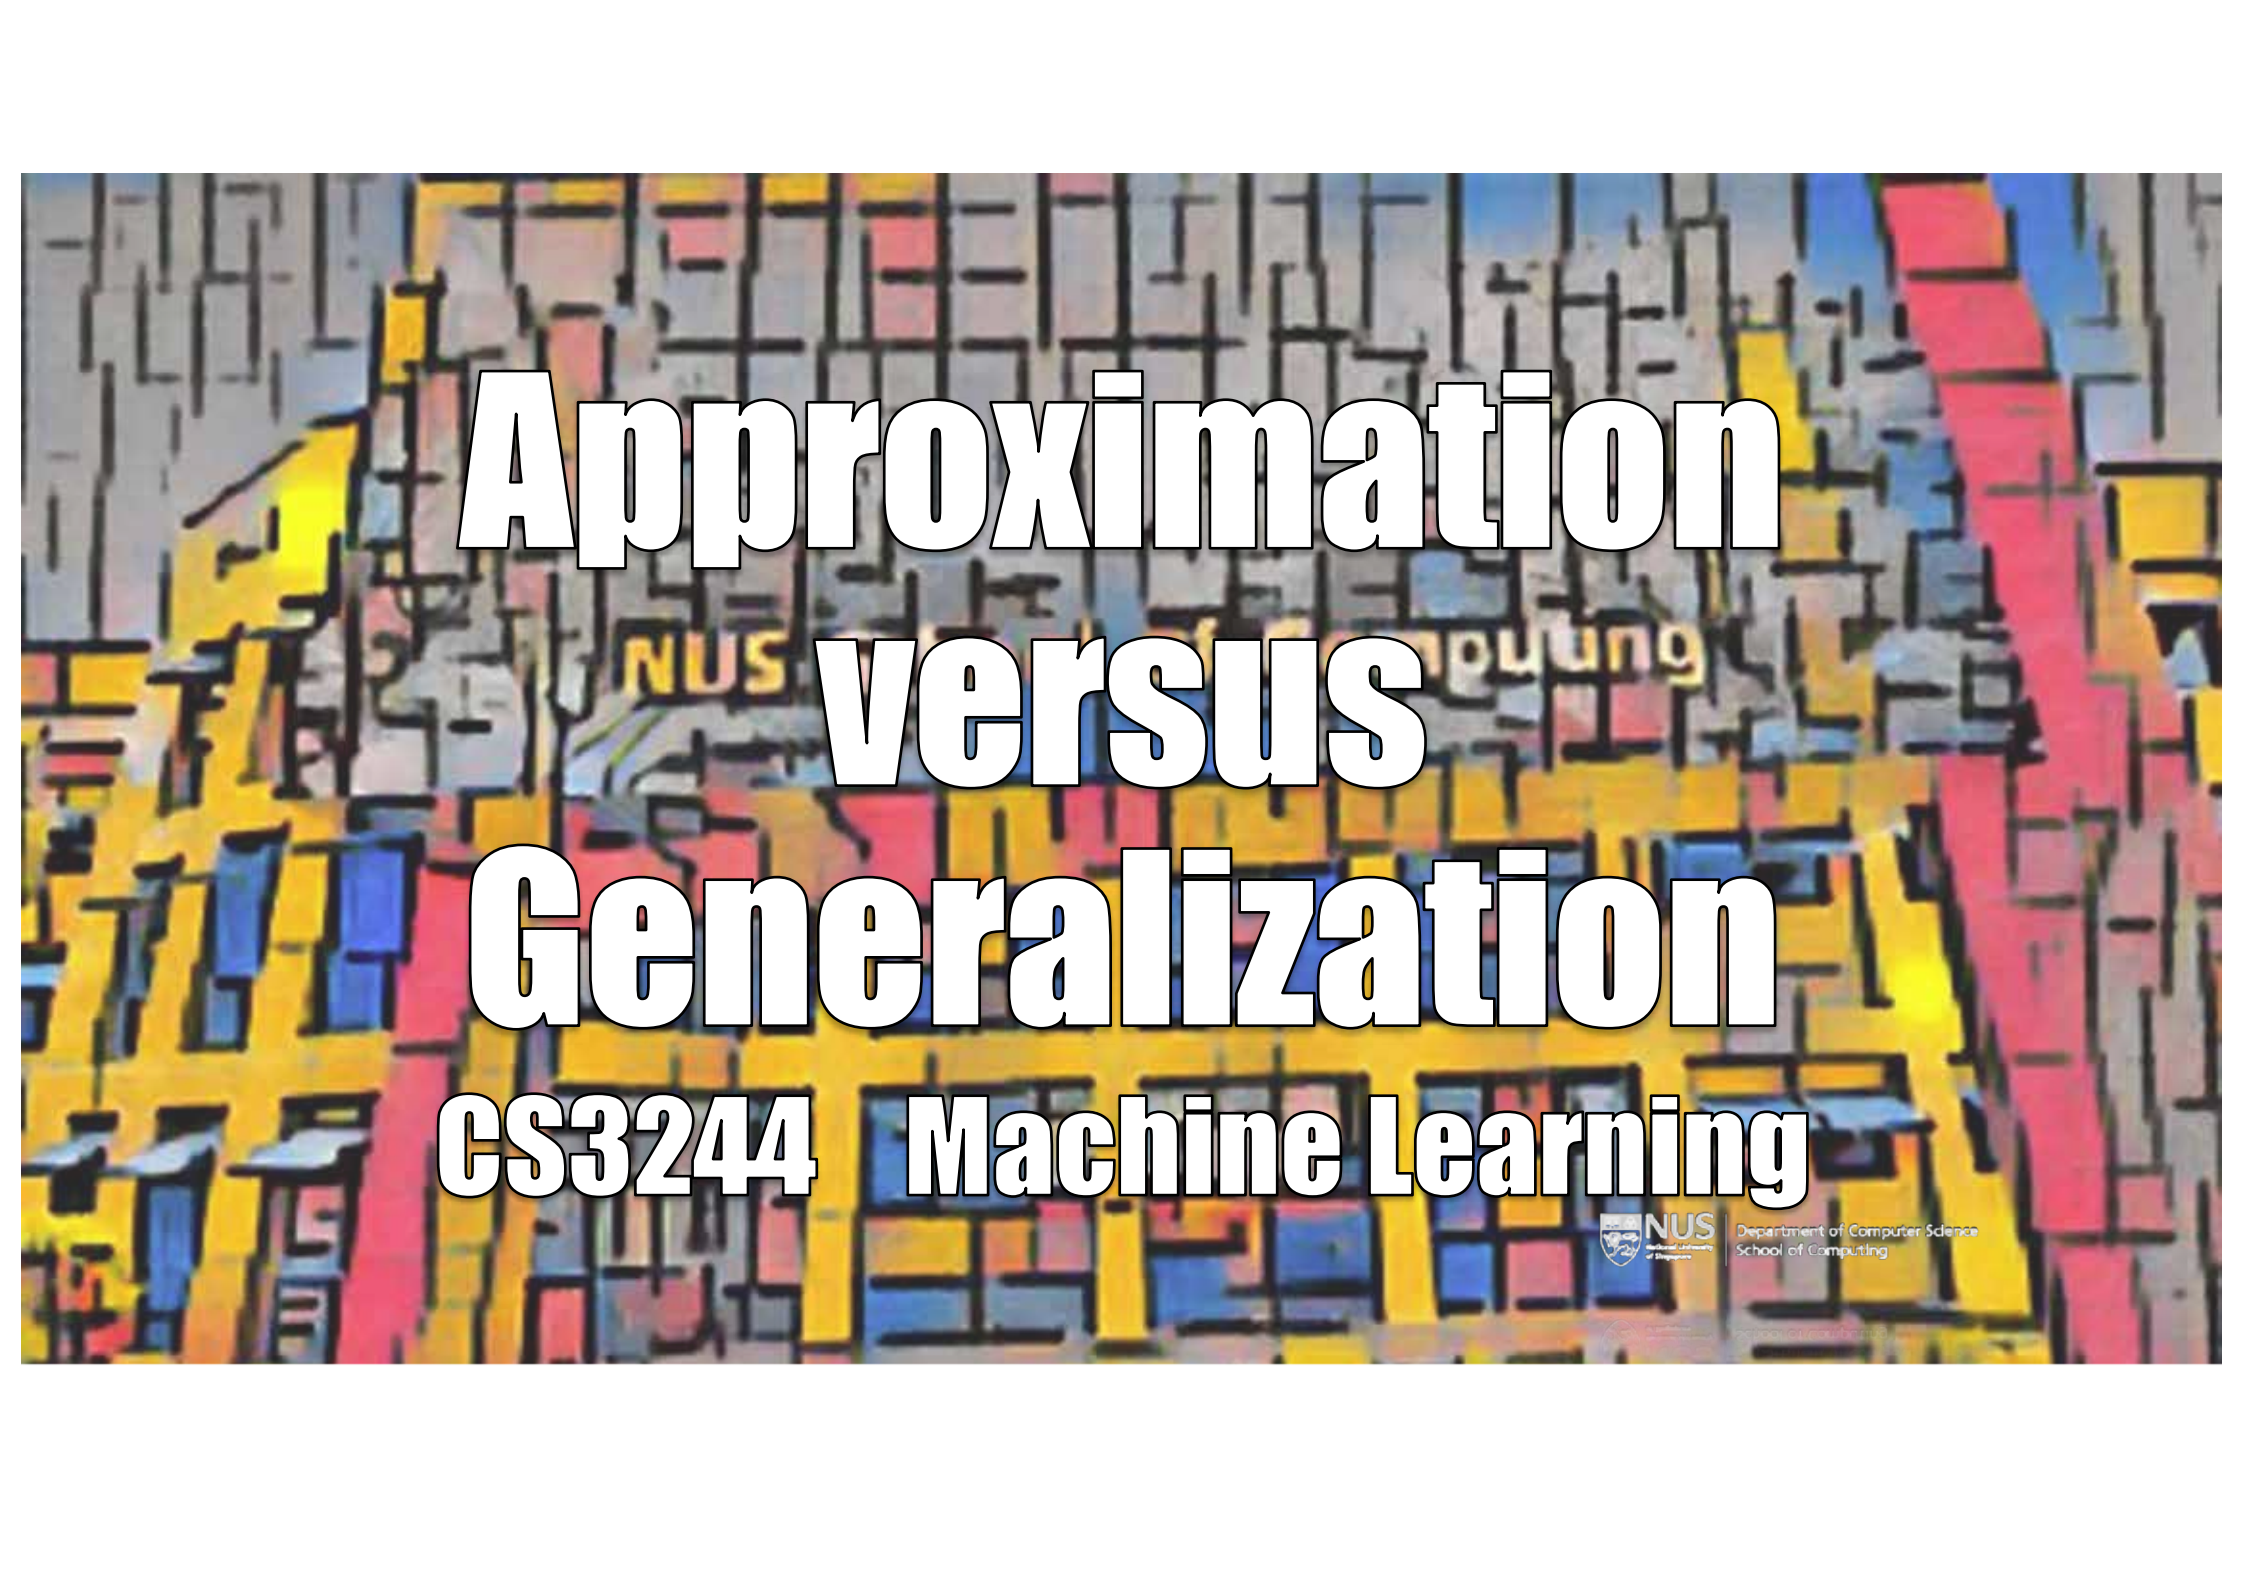


	2 nd Order	
L_{train}	0.050	
L_{test}	0.127	

Noiseless high-order target



	2 nd Order	
L_{train}	0.029	
L_{test}	0.120	



Approximation versus Generalization

CS3244 Machine Learning



Department of Computer Science
School of Computing

Approximation–generalization tradeoff



Small L_{test} : good approximation of f on unseen test data

More complex $\mathcal{H} \rightarrow$ better chance of **approximating** f

Less complex $\mathcal{H} \rightarrow$ better chance of **generalizing** on test data

Ideal $\mathcal{H} = \{f\}$

Good luck
with that!

Occam's Razor

CORE PRINCIPLES IN RESEARCH



OCCAM'S RAZOR

"WHEN FACED WITH TWO POSSIBLE EXPLANATIONS, THE SIMPLER OF THE TWO IS THE ONE MOST LIKELY TO BE TRUE."



OCCAM'S PROFESSOR

"WHEN FACED WITH TWO POSSIBLE WAYS OF DOING SOMETHING, THE MORE COMPLICATED ONE IS THE ONE YOUR PROFESSOR WILL MOST LIKELY ASK YOU TO DO."

WWW.PHDCOMICS.COM

JORGE CHAM © 2009

"Simpler is better"

Formalized as **Minimum Description Length (MDL)**

The shortest description of the data is the best model (related to **Entropy**)

Last Week's Assigned Task



Post a 1–2 sentence answer to the topic in your tutorial group:


#tg-xx

Describe kNN or decision trees with respect to variance.

In k NN

Model free: Data completely dictates the model


Higher k lowers the dependence of the model on a particular data point, makes model more robust and lowers variance




...In Decision Trees

Pruning discards detailed tests that may use criteria non-essential for classification in test data.

Before



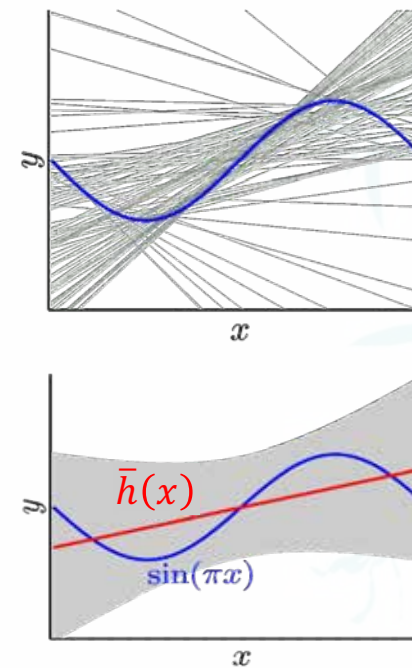

After




...In Linear Models

We just did this!

Each additional parameter θ_i adds a degree of freedom in the model.



...In SVM



The SVM determines its separating hyperplane by virtue of its **support vectors**.

of support vectors \propto complexity of the model

... in Ensembles

We have T reports h_1, h_2, \dots, h_T predicting whether a stock will go up as $h(x)$.

We can:

1. Select the most trustworthy of them based on their usual performance

Training performance: $h(x) = h_{t^*}(x)$ with $t^* = \operatorname{argmin}_{t \in \{1, 2, \dots, T\}} L_{\text{train}}(h_t)$

2. Let each report have a vote

Uniform Vote: $h(x) = \operatorname{sign}(\sum_{t=1}^T 1 \cdot h_t(x))$

3. Weight the reports non-uniformly

Weighted Vote: $h(x) = \operatorname{sign}(\sum_{t=1}^T \alpha_t \cdot h_t(x))$ where $\alpha_t \geq 0$.

4. Combine the predictions conditionally

This is decision trees, let's finish it up now!

Key: component hypotheses are
Diverse.

Ensembling diverse h
generalizes better.

 **fact:** Decision trees have a
native version: **Random forests.**



Bias and Variance, Analytically

CS3244 Machine Learning



Department of Computer Science
School of Computing

Quantifying the tradeoff

Decomposing L_{test} into

1. How well can \mathcal{H} approximate f overall
2. How well we can zoom in on a good $h_\theta \in \mathcal{H}$

Applies to **real-valued** targets and uses **squared error**

There is an equivalent binary version through
the lens of VC analyses (related to SVM)

Recap: Expected Value

We'll need to reason about values of data, irrespective of particular samples, so we'll need to work with **expected values of random variable**.

Expected Value: Average over all possible outcome of a random variable X .

$$\mathbb{E}[X] = \sum_{x \in X} x \Pr[X = x]$$

Not conditional test notation
(not Iverson brackets)

Example

Let X be the outcome of a fair dice roll.

Then the expected value is

$$\mathbb{E}[X] = \frac{1}{6}(1 + 2 + 3 + 4 + 5 + 6) = 3.5.$$

Variance

Intuition

Quantify how much deviation from expected value.

$$\text{Var}[x] = \mathbb{E}[(x - \mathbb{E}[x])^2]$$

Example

Let x be the outcome of a fair dice roll. Then the variance is

$$\text{Var}[x] = \frac{1}{6} \left((1 - 3.5)^2 + \dots + (6 - 3.5)^2 \right) = \frac{35}{12}.$$

Some Terminology



- $\mathbb{E}_x[z(\textcolor{red}{x})] \equiv$ Expected value of $z()$, given the distribution of values of $\textcolor{red}{x}$.
- $h_{\textcolor{blue}{\mathcal{D}}} \equiv$ Hypothesis of learner when learning from Dataset $\textcolor{blue}{\mathcal{D}}$.
- $\mathbb{E}_{\textcolor{blue}{\mathcal{D}}}[z()] \equiv$ Expected value of $z()$, given distribution of possible Datasets $\textcolor{blue}{\mathcal{D}}$.

Start with L_{test}

$$L_{test}(h_{\mathcal{D}}) = \mathbb{E}_x[(h_{\mathcal{D}}(x) - f(x))^2]$$



$h(x)$ depends on
the specific dataset \mathcal{D} .

Start with L_{test}

$$L_{test}(h_{\mathcal{D}}) = \mathbb{E}_x[(h_{\mathcal{D}}(x) - f(x))^2]$$

$$\mathbb{E}_{\mathcal{D}}[L_{test}(h_{\mathcal{D}})]$$

$$= \mathbb{E}_{\mathcal{D}}[\mathbb{E}_x[(h_{\mathcal{D}}(x) - f(x))^2]]$$

$$= \mathbb{E}_x[\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - f(x))^2]]$$

Focus on $\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - f(x))^2]$, we'll deal with the outer term $\mathbb{E}_x[\dots]$ later.



$h(x)$ depends on
the specific dataset \mathcal{D} .

Generalizing over all
 \mathcal{D} with same m

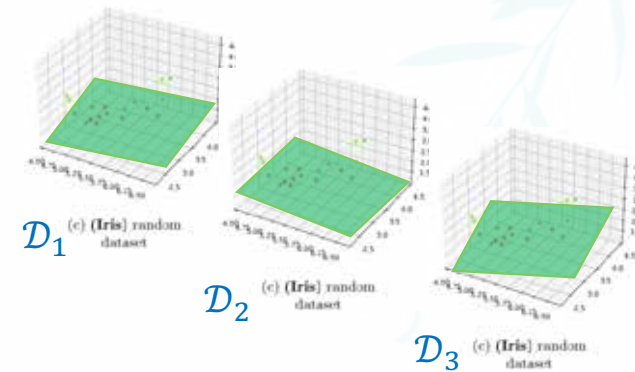

Swap ok, as integrand
is strictly non-negative

The average hypothesis

To evaluate $\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(\mathbf{x}) - f(\mathbf{x}))^2]$,
we define the ‘average’ hypothesis $\bar{h}(\mathbf{x}) = \mathbb{E}_{\mathcal{D}}[h_{\mathcal{D}}(\mathbf{x})]$

Imagine many, many data sets $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_K$
drawn:

$$\bar{h}(\mathbf{x}) \approx \frac{1}{K} \sum_{k=1}^K h_{\mathcal{D}_k}(\mathbf{x})$$



Using $\bar{h}(x)$

$$\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - f(x))^2] =$$

Using the average hypothesis $\bar{h}(x)$



$$\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - f(x))^2] =$$

$$= \mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - \bar{h}(x) + \bar{h}(x) - f(x))^2]$$

Add $-\bar{h}(x) + \bar{h}(x)$.
Associate first two and
second two terms.

$$= \mathbb{E}_{\mathcal{D}} \left[\underbrace{(h_{\mathcal{D}}(x) - \bar{h}(x))^2}_0 + (\bar{h}(x) - f(x))^2 + 2 \underbrace{(h_{\mathcal{D}}(x) - \bar{h}(x))}_{\text{const w.r.t. } \mathcal{D}} \underbrace{(\bar{h}(x) - f(x))}_{\text{const w.r.t. } \mathcal{D}} \right]$$

Expand out, cross terms
drop as first part of term is 0.

$$= \mathbb{E}_{\mathcal{D}} [(h_{\mathcal{D}}(x) - \bar{h}(x))^2] + (\bar{h}(x) - f(x))^2$$

2nd term is constant with
respect to \mathcal{D}

Bias and Variance



$$\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(\mathbf{x}) - f(\mathbf{x}))^2] = \underbrace{\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(\mathbf{x}) - \bar{h}(\mathbf{x}))^2]}_{\text{var}} + \underbrace{(\bar{h}(\mathbf{x}) - f(\mathbf{x}))^2}_{\text{bias}}$$

Therefore,

$$\begin{aligned}\mathbb{E}_{\mathcal{D}}[L_{test}(h_{\mathcal{D}})] &= \mathbb{E}_{\mathbf{x}}[\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(\mathbf{x}) - f(\mathbf{x}))^2]] \\ &= \mathbb{E}_{\mathbf{x}}[\text{bias}(\mathbf{x}) + \text{var}(\mathbf{x})] \\ &\equiv \text{bias} + \text{var}\end{aligned}$$





The tradeoff



$$\text{Bias}[h(x)] = |\bar{h}(x) - f(x)|.$$

$$\text{Var}[h(x)] = \mathbb{E} \left[\left(h(x) - \bar{h}(x) \right)^2 \right]$$





Overfit Measure

CS3244 Machine Learning





Department of Computer Science
School of Computing

10th order polynomial with noise




Two learners **Overfitting** and **Restrained**:



Both know the target is 10th order,
you get 15 data points.

- 1  chooses \mathcal{H}_{10}
- 2  chooses \mathcal{H}_2

Your Turn:
Which fits
better?



When is \mathcal{H}_2 better than \mathcal{H}_{10} ?



Overfitting (grey region): $L_{test}(\mathcal{H}_{10}) > L_{test}(\mathcal{H}_2)$


Without noise, with complex f

The two learners \mathcal{H}_{10} and \mathcal{H}_2

They know there is no noise


Is there really no noise?

Learning a 50th
order target



A detailed experiment

Impact of **noise level** and **target complexity**



$$y = f(x) + \epsilon(x) = \underbrace{\sum_{q=0}^{Q_f} \alpha_q x^q}_* + \underbrace{\epsilon(x)}_{\sigma^2}$$

Noise level: σ^2

Target complexity: Q_f


Data set size: m

The overfit measure

We fit the data set $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$ using our 2 models:

\mathcal{H}_2 : 2nd order polynomials

\mathcal{H}_{10} : 10th order polynomials




Compare out-of-sample errors of $h_2 \in \mathcal{H}_2$ and $h_{10} \in \mathcal{H}_{10}$

Overfit measure

$$L_{test}(h_{10}) - L_{test}(h_2)$$

Overfit measure: $L_{test}(h_{10}) - L_{test}(h_2)$



Number of data points	↑	Overfitting	↓
Stochastic Noise	↑	Overfitting	↑
Deterministic Noise	↑	Overfitting	↑



The Role of Noise

CS3244 Machine Learning



Department of Computer Science
School of Computing

Noise



That part of y that
we **cannot** model

It has two sources ...


Stochastic Noise: Data Error

We would like to learn f



$$y_j = f(\mathbf{x}_j)$$

Unfortunately, we actually observe \mathcal{D}

$$y_j = f(\mathbf{x}_j) + \text{stochastic_noise}$$



Stochastic Noise
Fluctuations that
we cannot model



Deterministic Noise: Model Error


We would like to learn from \bigcirc

$$y_j = \bar{h}(\mathbf{x}_j)$$

Unfortunately, we only observe \mathcal{D}



$$y_j = f(\mathbf{x}_j)$$

$$= \bar{h}(\mathbf{x}_j) + \text{deterministic_noise}$$




Deterministic Noise

The part of f we
lack the capacity
to model



Both sources of noise hurt learning

Stochastic Noise




Source: random measurement errors

Re-measure y : Stochastic noise changes

Change \mathcal{H} : Stochastic noise still the same

Deterministic Noise



Source: learner's \mathcal{H} cannot model f

Re-measure y : Deterministic noise the same

Change \mathcal{H} : Deterministic noise changes

We have a single \mathcal{D} and fixed \mathcal{H} so we cannot distinguish between either; with finite m , \mathcal{H} will try to fit the noise.

Noise and Bias-Variance



Recall the decomposition:

$$\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - f(x))^2] = \underbrace{\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - \bar{h}(x))^2]}_{\text{variance}} + \underbrace{(\bar{h}(x) - f(x))^2}_{\text{bias}}$$

What if f is a **noisy** target?

$$y = f(x) + \epsilon(x)$$

$$\mathbb{E}[\epsilon(x)] = 0$$

A noise term

$$\mathbb{E}_{\mathcal{D}, \epsilon}[(h_{\mathcal{D}}(x) - y)^2] =$$

A noise term



$$\mathbb{E}_{\mathcal{D}, \epsilon}[(h_{\mathcal{D}}(x) - y)^2] =$$

$$= \mathbb{E}_{\mathcal{D}, \epsilon}[(h_{\mathcal{D}}(x) - f(x) - \epsilon(x))^2]$$

Expand observed y

$$= \mathbb{E}_{\mathcal{D}, \epsilon} \left[\left(\underbrace{h_{\mathcal{D}}(x) - \bar{h}(x)}_{\textcircled{1}} + \underbrace{\bar{h}(x) - f(x)}_{\textcircled{2}} - \underbrace{\epsilon(x)}_{\textcircled{3}} \right)^2 \right]$$

Associate terms


$$= \mathbb{E}_{\mathcal{D}, \epsilon} \left[\left(h_{\mathcal{D}}(x) - \bar{h}(x) \right)^2 + \left(\bar{h}(x) - f(x) \right)^2 + (\epsilon(x))^2 + \text{cross terms} \right]$$

Additional cross
terms disappear as
 $\mathbb{E}(\epsilon(x)) = 0$


Actually, two noise terms

$$\mathbb{E}_{\mathcal{D}}[(h_{\mathcal{D}}(x) - f(x))^2] =$$

$$\underbrace{\mathbb{E}_{\mathcal{D},x}[(h_{\mathcal{D}}(x) - \bar{h}(x))^2]}_{\text{variance}} + \underbrace{\mathbb{E}_x(\bar{h}(x) - f(x))^2}_{\text{bias}} + \underbrace{\mathbb{E}_{\epsilon,x}(\epsilon(x))^2}_{\sigma^2}$$



Deterministic Noise



Stochastic Noise



Wrapping up Week 05

CS3244 Machine Learning




Department of Computer Science
School of Computing

Summary

Bias-Variance Tradeoff


Bias: How well \mathcal{H} can approximate f overall

Variance: How well we can zoom in on a good $h \in \mathcal{H}$



Match the 'model complexity' to the **data resources**, not to the target complexity

Overfitting: Fitting the data more than is warranted



Two causes: stochastic + deterministic noise

Bias \equiv deterministic noise

Noise causes overfitting

Overfitting is the disease

Noise is the cause

Learning is led astray by fitting the noise more than the signal

Two Cures:

1. **Regularization**: Restrain the model
2. **Validation**: Reality check by peeking at (the bottom line)

Outlook for next week

Assigned Task (due before next Mon)



Read the post <https://www.kaggle.com/alexisbcook/cross-validation>
(10 mins)

Post a 1–2 sentence answer to the topic in your tutorial group: **#tg-xx**

How does cross validation relate to variance?

[There's an optional exercise with this course page, using random forest and MAE, you're welcomed to do it if you'd like]